

Interaction-Partitioned Topic Models (IPTM) using a Point Process Approach

Bomin Kim¹, Bruce Desmarais¹, and Hanna Wallach^{2,3}

¹Pennsylvania State University

²Microsoft Research NYC

³University of Massachusetts Amherst

November 13, 2016

1 Ideas

Current CPME model does not involve any of temporal component, which plays a key role in email interactions. Intuitively, past interaction behaviors significantly influence future ones; for example, if an actor i sent an email to actor j , then j is highly likely to send an email back to i as a response (i.e. reciprocity). Moreover, the recency and frequency of past interactions can also be considered to effectively predict future interactions. Thus, as an exploratory data analysis, point process model for directional interaction is applied to the North Carolina email data. Starting from the existing framework focused on the analysis of content-partitioned subnetworks, I would suggest an extended approach to analyze the data using the timestamps in the email, aiming to develop a joint dynamic or longitudinal model of text-valued ties.

CPME model is a Bayesian framework using two well-known methods: Latent Dirichlet Allocation (LDA) and Latent Space Model (LSM). Basically, existence of edge depends on topic assignment k (LDA) and its corresponding interaction pattern c . Each topic $k = 1, \dots, K$ has one interaction pattern $c=1, \dots, C$, and each interaction pattern posits unique latent space (LSM), thus generating $A \times A$ matrix of probabilities $P^{(c)}$ that a message author a will include recipient r on the message, given that it is about a topic in cluster c . Incorporating point process approach, now assume that under each interaction pattern, we have $A \times A$ matrix of stochastic intensities at time t , $\lambda^{(c)}(t)$, which depends on $\mathbf{x}_t^{(c)}(i, j)$, the history of interaction between the sender and receiver corresponding to the interaction pattern c . We will refer this as interaction-partitioned topic models (IPTM).

2 IPTM Model

In this section, we introduce multiplicative Cox regression model for the edge formation process in a longitudinal communication network. For concreteness, we frame our discussion of this model in terms of email data, although it is generally applicable to any similarly-structured communication data.

2.1 Point Process Framework

A single email, indexed by d , is represented by a set of tokens $w^{(d)} = \{w_m^{(d)}\}_{m=1}^{M^{(d)}}$ that comprise the text of that email, an integer $i^{(d)} \in \{1, \dots, A\}$ indicating the identity of that email's sender, an integer $j^{(d)} \in \{1, \dots, A\}$ indicating the identity of that email's receiver, and an integer $t^{(d)} \in [0, T]$ indicating the (unix time-based) timestamp of that email. To capture the relationship between the interaction patterns expressed in an email and that email's recipients, documents that share the interaction pattern c are associated with an $A \times A$ matrix of $\lambda^{(c)}(t) = \{\{\lambda_{ij}^{(c)}(t)\}_{i=1}^A\}_{j=1}^A$, the stochastic

intensity where $\lambda_{ij}^{(c)}(t)dt = P\{\text{for interaction pattern } c, i \rightarrow j \text{ occurs in time interval } [t, t + dt)\}$. We will model the counting process $\mathbf{N}^{(d|c)}(t)$ through $\boldsymbol{\lambda}^{(c)}(t)$ using a version of the Cox proportional intensity model, where $N_{ij}^{(d|c)}(t)$ denotes the number of edges (emails) for document d from actor i to actor j up to time t (from the starting point 0) given that the document corresponds to interaction pattern c . Since this counting process \mathbf{N} is document-based, each element is either 0 or 1, and only one element of the matrix is 1 while all the rests are 0 (assuming no multicast).

Combining the individual counting processes of all potential edges, $\mathbf{N}^{(d|c)}(t)$ is the multivariate counting process with $\mathbf{N}^{(d|c)}(t) = (N_{ij}^{(d|c)}(t) : i, j \in 1, \dots, A, i \neq j)$. Here we make no assumption about the independence of individual edge counting process. As in Vu et al. (2011), we model the multivariate counting process via Doob-Meyer decomposition:

$$\mathbf{N}^{(d|c)}(t) = \int_0^t \boldsymbol{\lambda}^{(c)}(s)ds + \mathbf{M}(t) \quad (1)$$

where essentially $\boldsymbol{\lambda}^{(c)}(t)$ and $\mathbf{M}(t)$ may be viewed as the (deterministic) signal and (martingale) noise, respectively.

Following the multiplicative Cox model of the intensity process $\boldsymbol{\lambda}^{(c)}(t)$ given $\mathbf{H}_{t-}^{(c)}$, the entire past of the network corresponding to the interaction pattern c up to but not including time t , we consider for each potential directed edge (i, j) the intensity forms:

$$\lambda_{ij}^{(c)}(t|\mathbf{H}_{t-}^{(c)}) = \lambda_0 \cdot \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_t^{(c)}(i, j)\} \cdot 1\{j \in \mathcal{A}^{(c)}\} \quad (2)$$

where λ_0 is the common baseline hazards for the overall interaction, $\boldsymbol{\beta}^{(c)}$ is an unknown vector of coefficients in \mathbf{R}^p , $\mathbf{x}_t^{(c)}(i, j)$ is a vector of p statistics for directed edge (i, j) constructed based on $\mathbf{H}_{t-}^{(c)}$, and $\mathcal{A}^{(c)}$ is the predictable receiver set of sender i corresponding to the interaction pattern c within the set of all possible actors \mathcal{A} . Equivalently, by fixing $\lambda_0 = 1$, we can rewrite (2):

$$\lambda_{ij}^{(c)}(t|\mathbf{H}_{t-}^{(c)}) = \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_t^{*(c)}(i, j)\} \cdot 1\{j \in \mathcal{A}^{(c)}\} \quad (3)$$

where the first element of $\boldsymbol{\beta}^{(c)}$ corresponds to the deviation from λ_0 , by setting $\mathbf{x}_t^{*(c)}(i, j) = (\mathbf{1}, \mathbf{x}_t^{(c)}(i, j))$.

Based on the framework illustrated so far, the likelihood we will use for inference procedure is that of Perry and Wolfe (2013). For each type of interaction pattern $c = 1, \dots, C$, estimation for $\boldsymbol{\beta}^{(c)}$ proceeds by maximizing the so-called partial likelihood of Cox (1992):

$$PL_t(\boldsymbol{\beta}^{(c)}) = \prod_{d:c^{(d)}=c} \frac{\exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)})\}}{\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c)}(i^{(d)}, j)\}}, \quad (4)$$

where $t^{(d)}$, $i^{(d)}$, and $j^{(d)}$ are the time, sender, and receiver of the d th document. For computational efficiency, we will use the log-partial likelihood:

$$\log PL_t(\boldsymbol{\beta}^{(c)}) = \sum_{d:c^{(d)}=c} \left\{ \boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)}) - \log \left[\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c)}(i^{(d)}, j)\} \right] \right\}. \quad (5)$$

2.2 Generative Process

The generative process of this model follows the topic model (LDA) of Blei et al. (2003) and the author-topic model of Rosen-Zvi et al. (2004). Same as LDA, documents are represented as random mixtures over latent topics, where each topic is characterized by a distribution over words. However, one crucial difference is that each document is connected to one type of interaction pattern, and the topic distributions vary depending on the assigned interaction pattern.

Conditioned on the interaction pattern and their distributions over topics, the process by which a document is generated can be summarized as follows: first, an interaction pattern is chosen by multinomial for each document; next, a topic is sampled for each word from the distribution over topics associated with the interaction pattern of the document; finally, words themselves are sampled from the distribution over words associated with each topic. At the same time, the unique sender-recipient pair of the document is determined by the rate of intensities associated with the interaction pattern and history of interactions until the time the document is written. Below are the detailed generative process for each document in a corpus D and its plate notation (Figure 1), and Table 1 summarizes the notations used in this paper:

1. $\phi^{(k)} \sim \text{Dir}(\delta, \mathbf{n})$ [See Algorithm 1]
 - A “topic” k is characterized by a discrete distribution over V word types with probability vector $\phi^{(k)}$. A symmetric Dirichlet prior with concentration parameter δ is placed.
2. For each of the C interaction patterns [See Algorithm 2]:
 - (a) $\beta^{(c)} \sim \text{Normal}(\mathbf{0}, \sigma^2 I_P)$
 - The vector of coefficients depends on the interaction pattern c . This means that there is variation in the degree of influence from the network statistics.
 - (b) Update $\mathbf{x}_t^{*(c)}(i, j)$
 - Corpus are partitioned according to the assignment of interaction patterns, and the dynamic network statistics are calculated based on the documents of the same interaction pattern.
 - (c) Using $\beta^{(c)}$ in (a), update $\lambda^{(c)}(t)$
 - Use the equation $\lambda_{ij}^{(c)}(t) = \exp\left\{\beta^{(c)T} \mathbf{x}_t^{*(c)}(i, j)\right\} \cdot 1\{j \in \mathcal{A}^{(c)}\}$ for all $i \in \mathcal{A}, j \in \mathcal{A}, i \neq j$.
 - (c) $\theta^{(c)} \sim \text{Dir}(\alpha, \mathbf{m})$
 - Each email has a discrete distribution over topics $\theta^{(c)}$, since the topic proportions for documents in the same cluster are drawn from the same distribution. The Dirichlet parameters α and \mathbf{m} may or may not vary by interaction patterns.
3. For each of the D documents [See Algorithm 3]:
 - (a) $c^{(d)} \sim \text{Multinomial}(\gamma)$
 - Each document d is associated with one “interaction pattern” among C different types, with parameter γ . Here, we assign the prior for the multinomial parameter $\gamma \sim \text{Dir}(\eta, \mathbf{l})$
 - (b) $\mathbf{N}^{(d|c^{(d)})}(t^{(d)}) \sim \text{CP}(\lambda^{(c^{(d)})}(t^{(d)}))$
 - The actual update of the counting process $\mathbf{N}^{(d|c^{(d)})}(t)$ of the email d is $N_{i^{(d)}j^{(d)}}^{(d|c^{(d)})}(t^{(d)}) = 1$ and the rest $N_{(i,j) \neq (i^{(d)}, j^{(d)})}^{(d|c^{(d)})}(t^{(d)}) = 0$.
4. For each of the M words [See Algorithm 4]:
 - (a) $z_m^{(d)} \sim \text{Multinomial}(\theta^{(c^{(d)})})$
 - (b) $w_m^{(d)} \sim \text{Multinomial}(\phi^{(z_m^{(d)})})$

Algorithm 1 Topic Word Distributions

for $k=1$ **to** K **do**
 | draw $\phi^{(k)} \sim \text{Dir}(\delta, \mathbf{n})$
end

Algorithm 2 Interaction Patterns

```
for  $c=1$  to  $C$  do
  draw  $\beta^{(c)} \sim \text{Normal}(\mathbf{0}, \sigma^2 I_P)$ 
  set  $\mathbf{x}_t^{*(c)}(i, j)$  according to Section 2.3
  for  $i=1$  to  $A$  do
    for  $j=1$  to  $A$  do
      if  $i \neq j$  then
        set  $\lambda_{ij}^{(c)}(t) = \exp\{\beta^{(c)T} \mathbf{x}_t^{*(c)}(i, j)\} \cdot 1\{j \in \mathcal{A}^{(c)}\}$ 
      end
    else
      set  $\lambda_{ij}^{(c)}(t) = 0$ 
    end
  end
end
draw  $\theta^{(c)} \sim \text{Dir}(\alpha, \mathbf{m})$ 
end
```

Algorithm 3 Document-Interaction Pattern Assignments

```
for  $d=1$  to  $D$  do
  draw  $c^{(d)} \sim \text{Multinomial}(\gamma)$ 
  draw  $\mathbf{N}^{(d|c^{(d)})}(t^{(d)}) \sim \text{CP}(\lambda^{(c^{(d)})}(t^{(d)}))$ 
end
```

Algorithm 4 Tokens

```
for  $d=1$  to  $D$  do
  set  $M^{(d)}$  = the number of words in document  $d$ 
  for  $m=1$  to  $M^{(d)}$  do
    draw  $z_m^{(d)} \sim \text{Multinomial}(\theta^{(c^{(d)})})$ 
    draw  $w_m^{(d)} \sim \text{Multinomial}(\phi^{(z_m^{(d)})})$ 
  end
end
```

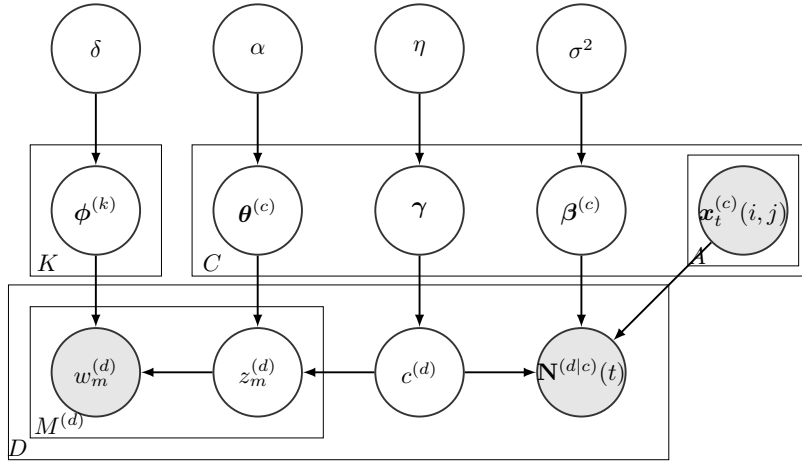


Figure 1: Plate notation of IPTM

Authors of the corpus	\mathcal{A}	Set
Authors of the corpus given interaction pattern c	$\mathcal{A}^{(c)}$	Set
Number of authors	A	Scalar
Number of documents	D	Scalar
Number of words in the d^{th} document	$M^{(d)}$	Scalar
Number of topics	K	Scalar
Vocabulary size	W	Scalar
Number of interaction patterns	C	Scalar
Number of words assigned to interaction pattern and topic	M^{CK}	Scalar
Number of words assigned to word and topic	M^{WK}	Scalar
Interaction pattern of the d^{th} document	$c^{(d)}$	Scalar
Time of the d^{th} document	$t^{(d)}$	Scalar
Words in the d^{th} document	$\mathbf{w}^{(d)}$	$M^{(d)}$ -dimensional vector
m^{th} word in the d^{th} document	$w_m^{(d)}$	m^{th} component of $\mathbf{w}^{(d)}$
Topic assignments in the d^{th} document	$\mathbf{z}^{(d)}$	$M^{(d)}$ -dimensional vector
Topic assignments for m^{th} word in the d^{th} document	$z_m^{(d)}$	m^{th} component of $\mathbf{z}^{(d)}$
Dirichlet concentration prior	α	Scalar
Dirichlet base prior	\mathbf{m}	K -dimensional vector
Dirichlet concentration prior	δ	Scalar
Dirichlet base prior	\mathbf{n}	W -dimensional vector
Dirichlet concentration prior	η	Scalar
Dirichlet base prior	\mathbf{l}	C -dimensional vector
Multinomial prior	γ	C -dimensional vector
Variance of Normal prior	σ^2	Scalar
Probabilities of the words given topics	Φ	$W \times K$ matrix
Probabilities of the words given topic k	$\phi^{(k)}$	W -dimensional vector
Probabilities of the topics given interaction patterns	Θ	$K \times C$ matrix
Probabilities of the topics given interaction pattern c	$\theta^{(c)}$	K -dimensional vector
Coefficient of the intensity process given interaction pattern c	$\beta^{(c)}$	p -dimensional vector
Network statistics for directed edge (i, j) given interaction pattern c	$\mathbf{x}_t^{(c)}(i, j)$	p -dimensional vector
Counting process in the d^{th} document given interaction pattern	$\mathbf{N}^{(d c)}(t)$	$A \times A$ matrix

Table 1: Symbols associated with IPTM, as used in this work

2.3 Dynamic covariates to measure network effects

The network statistics $\mathbf{x}_t^{(c)}(i, j)$ of Equation (2), corresponding to the ordered pair (i, j) , can be time-invariant (such as gender) or time-dependent (such as the number of two-paths from i to j just before time t). Since time-invariant covariates can be easily specified in various manners (e. g. homophily or group-level effects), here we only consider specification of dynamic covariates.

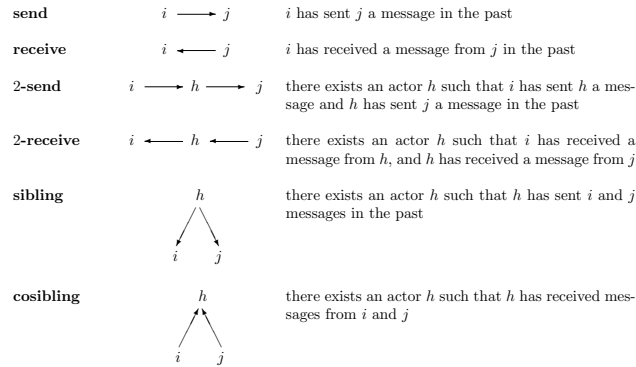


Fig. 3. Dynamic covariates to measure network effects

Following Perry and Wolfe (2013) (refer to Fig.3 of Perry and Wolfe (2013) attached above), we use 6 effects as components of $\mathbf{x}_t^{(c)}(i, j)$. The first two behaviors (send and receive) are dyadic, involving exactly two actors, while the last four (2-send, 2-receive, sibling, and cosibling) are triadic, involving exactly three actors. In addition, we include intercept term and use $\mathbf{x}_t^{*(c)}(i, j)$ so that we can estimate the baseline intensities at the same time. However, one different thing from the existing specification

is that we define the effects not to be based on finite sub-interval, which require large number of dimension. Instead, we create a single statistic for each effect by incorporating the recency of event into the statistic itself.

$$0. \text{ intercept}_t(i, j) = 1$$

$$1. \text{ send}_t(i, j) = \sum_{d:t^{(d)} < t} I\{i \rightarrow j\} \cdot g(t - t^{(d)})$$

$$2. \text{ receive}_t(i, j) = \sum_{d:t^{(d)} < t} I\{j \rightarrow i\} \cdot g(t - t^{(d)})$$

$$3. \text{ 2-send}_t(i, j) = \sum_{h \neq i, j} \left(\sum_{d:t^{(d)} < t} I\{i \rightarrow h\} \cdot g(t - t^{(d)}) \right) \left(\sum_{d:t^{(d)} < t} I\{h \rightarrow j\} \cdot g(t - t^{(d)}) \right)$$

$$4. \text{ 2-receive}_t(i, j) = \sum_{h \neq i, j} \left(\sum_{d:t^{(d)} < t} I\{h \rightarrow i\} \cdot g(t - t^{(d)}) \right) \left(\sum_{d:t^{(d)} < t} I\{j \rightarrow h\} \cdot g(t - t^{(d)}) \right)$$

$$5. \text{ sibling}_t(i, j) = \sum_{h \neq i, j} \left(\sum_{d:t^{(d)} < t} I\{h \rightarrow i\} \cdot g(t - t^{(d)}) \right) \left(\sum_{d:t^{(d)} < t} I\{h \rightarrow j\} \cdot g(t - t^{(d)}) \right)$$

$$6. \text{ cosibling}_t(i, j) = \sum_{h \neq i, j} \left(\sum_{d:t^{(d)} < t} I\{i \rightarrow h\} \cdot g(t - t^{(d)}) \right) \left(\sum_{d:t^{(d)} < t} I\{j \rightarrow h\} \cdot g(t - t^{(d)}) \right)$$

Here, $g(t - t^{(d)})$ reflects the difference between current time t and the timestamp of previous email $t^{(d)}$, thus measuring the recency. Inspired by the self-exciting Hawkes process, which is often used to model the temporal effect of email data, we can take the exponential kernel $g(t - t^{(d)}) = \lambda e^{-\lambda(t - t^{(d)})}$ where λ is the parameter of speed at which sender replies to emails, with larger values indicating faster response times. Indeed, λ^{-1} is the expected number of hours it takes to reply to a typical email. For simplicity, in our simulation we fixed $\lambda = 0.05$ and multiply 20 in front of exponential function to ensure $g = 1$ when $t = t^{(d)}$ (i.e. $g(t - t^{(d)}) = 20 \times 0.05 e^{-0.05(t - t^{(d)})}$), but this setup may vary based on the nature of document.

3 Inference

The inference for IPTM is similar to that of CPME. In this case, what we actually observe are the tokens $\mathcal{W} = \{\mathbf{w}^{(d)}\}_{d=1}^D$ and the sender, recipient, and timestamps of the email in the form of the counting process $\mathcal{N} = \{\mathbf{N}^{(d)}(t^{(d)})\}_{d=1}^D$. Next, $\mathcal{X} = \{\mathbf{x}_{t^{(d)}}^{(c)}(i, j)\}_{d=1}^D$ is the metadata, and the latent variables are $\Phi = \{\phi^{(k)}\}_{k=1}^K$, $\Theta = \{\theta^{(c)}\}_{c=1}^C$, $\mathcal{Z} = \{\mathbf{z}^{(d)}\}_{d=1}^D$, $\mathcal{C} = \{c^{(d)}\}_{d=1}^D$, and $\mathcal{B} = \{\beta^{(c)}\}_{c=1}^C$.

Below is the the big joint distribution

$$\begin{aligned} & P(\Phi, \Theta, \mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \eta, \sigma^2) \\ &= P(\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \Phi, \Theta, \mathcal{X}, \gamma, \eta, \sigma^2) P(\Phi, \Theta | \delta, \mathbf{n}, \alpha, \mathbf{m}) \\ &= P(\mathcal{W} | \mathcal{Z}, \Phi) P(\mathcal{Z} | \Theta) P(\mathcal{N} | \mathcal{C}, \mathcal{X}, \mathcal{B}) P(\mathcal{B} | \mathcal{C}, \sigma^2) P(\Phi | \delta, \mathbf{n}) P(\Theta | \mathcal{C}, \alpha, \mathbf{m}) P(\mathcal{C} | \gamma) P(\gamma | \eta) \end{aligned} \quad (6)$$

Now we can integrate out Φ and Θ in latent Dirichlet allocation by applying Dirichlet-multinomial conjugacy as we did in CPME. See APPENDIX A for the detailed steps. After integration, we obtain below:

$$\propto P(\mathcal{W} | \mathcal{Z}) P(\mathcal{Z} | \mathcal{C}, \delta, \mathbf{n}, \alpha, \mathbf{m}) P(\mathcal{N} | \mathcal{C}, \mathcal{B}, \mathcal{X}) P(\mathcal{B} | \mathcal{C}, \sigma^2) P(\mathcal{C} | \gamma) \quad (7)$$

Then, we only have to perform inference over the remaining unobserved latent variables \mathcal{Z}, \mathcal{C} , and \mathcal{B} , using the equation below:

$$P(\mathcal{Z}, \mathcal{C}, \mathcal{B} | \mathcal{W}, \mathcal{N}, \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \eta, \sigma^2) \propto P(\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \eta, \sigma^2) \quad (8)$$

Either Gibbs sampling or Metropolis-Hastings algorithm is applied by sequentially resampling each latent variables from their respective conditional posterior.

3.1 Resampling \mathcal{C}

The first variable we are going to resample is the document-interaction pattern assignments, one document at a time. To obtain the Gibbs sampling equation, which is the posterior conditional probability for the interaction pattern \mathcal{C} for d^{th} document, i.e. $P(c^{(d)} = c | \mathcal{W}, \mathcal{Z}, \mathcal{C}_{\setminus d}, \mathcal{B}, \mathcal{N}, \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2)$. We can derive the equation as below:

$$\begin{aligned} P(c^{(d)} = c | \mathcal{W}, \mathcal{Z}, \mathcal{C}_{\setminus d}, \mathcal{B}, \mathcal{N}, \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2) \\ \propto P(c^{(d)} = c, \mathbf{w}^{(d)}, \mathbf{z}^{(d)}, \mathbf{N}^{(d)}(t^{(d)}) | \mathcal{W}_{\setminus d}, \mathcal{Z}_{\setminus d}, \mathcal{C}_{\setminus d}, \mathcal{B}, \mathcal{N}_{\setminus d}, \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2) \\ \propto P(c^{(d)} = c | \mathcal{C}_{\setminus d}, \gamma) P(\mathbf{N}^{(d)}(t^{(d)}) | c^{(d)} = c, \mathcal{C}_{\setminus d}, \mathcal{B}, \mathcal{N}_{\setminus d}, \mathcal{X}) P(\mathbf{w}^{(d)}, \mathbf{z}^{(d)} | c^{(d)} = c, \mathcal{W}_{\setminus d}, \mathcal{Z}_{\setminus d}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m}), \end{aligned} \quad (9)$$

where $P(c^{(d)} = c | \mathcal{C}_{\setminus d}, \gamma)$ comes from the multinomial prior γ and $P(\mathbf{N}^{(d)}(t^{(d)}) | c^{(d)} = c, \mathcal{C}_{\setminus d}, \mathcal{B}, \mathcal{N}_{\setminus d}, \mathcal{X})$ is the probability of observing a document with the sender, receiver, and time equal to $(i = i^{(d)}, j = j^{(d)}, t = t^{(d)})$, respectively, given a set of parameter values. We will replace this by the partial likelihood in Equation (4) (without the product term since resampling of c is document-specific). For the last term $P(\mathbf{w}^{(d)}, \mathbf{z}^{(d)} | c^{(d)} = c, \mathcal{W}_{\setminus d}, \mathcal{Z}_{\setminus d}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m})$, we will follow typical LDA approach.

Using Bayes' theorem (See APPENDIX B for conditional probability of the last term), we have

$$= [\gamma_c] \times \left[\frac{\exp\{\boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)})\}}{\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j)\}} \right] \times \left[\prod_{m=1}^{M^{(d)}} \frac{M_{cz_m^{(d)}, \setminus d, m}^{CK} + \alpha m_k}{\sum_{k=1}^K M_{ck, \setminus d, m}^{CK} + \alpha} \right], \quad (10)$$

where M_{ck}^{CK} is the number of times topic k shows up across the documents of the interaction pattern c . Furthermore, we can take the log of Equation (10) to avoid numerical issue from exponentiation and increase the speed of computation, which becomes:

$$\log(\gamma_c) + \left(\boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)}) - \log \left[\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j)\} \right] \right) + \sum_{m=1}^{M^{(d)}} \log \left(\frac{M_{cz_m^{(d)}, \setminus d, m}^{CK} + \alpha m_k}{\sum_{k=1}^K M_{ck, \setminus d, m}^{CK} + \alpha} \right). \quad (11)$$

3.2 Resampling \mathcal{Z}

Next, the new values of $z_m^{(d)}$ are sampled for all of the token topic assignments (one token at a time), using the conditional posterior probability of being topic k as we derived in APPENDIX B:

$$\begin{aligned} P(z_m^{(d)} = k | \mathcal{W}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}, \mathcal{B}, \mathcal{N}, \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2) \\ \propto P(z_m^{(d)} = k, w_m^{(d)} | \mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}, \delta, \mathbf{n}, \alpha, \mathbf{m}) \end{aligned} \quad (12)$$

where the subscript " $\setminus d, m$ " denotes the exclusion of position m in email d . In the last line of equation (10), it is the contribution of LDA, so similar to CPME we can write the conditional probability:

$$\propto (M_{c^{(d)}k, \setminus d, m}^{CK} + \alpha m_k) \cdot \frac{M_{w_m^{(d)}k, \setminus d, m}^{WK} + \delta n_w}{\sum_{w=1}^W M_{wk, \setminus d, m}^{WK} + \delta} \quad (13)$$

which is the well-known form of collapsed Gibbs sampling equation for LDA.

3.3 Resampling \mathcal{B}

Finally, we want to update the interaction pattern parameter $\boldsymbol{\beta}^{(c)}$, one interaction pattern at a time. For this, we will use the Metropolis-Hastings algorithm with a proposal density Q being the multivariate Gaussian distribution, with variance δ_B^2 (proposal distribution variance parameters set by the user), centered on the current values of $\boldsymbol{\beta}^{(c)}$. Then we draw a proposal $\boldsymbol{\beta}'^{(c)}$ at each iteration.

Under symmetric proposal distribution (such as multivariate Gaussian), we cancel out Q-ratio and obtain the acceptance probability equal to:

$$\text{Acceptance Probability} = \begin{cases} \frac{P(\mathcal{B}'|\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{N}, \mathcal{X})}{P(\mathcal{B}|\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{N}, \mathcal{X})} & \text{if } < 1 \\ 1 & \text{else} \end{cases} \quad (14)$$

After factorization, we get

$$\begin{aligned} \frac{P(\mathcal{B}'|\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{N}, \mathcal{X})}{P(\mathcal{B}|\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{N}, \mathcal{X})} &= \frac{P(\mathcal{N}|\mathcal{B}', \mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{X})P(\mathcal{B}')}{P(\mathcal{N}|\mathcal{B}, \mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{X})P(\mathcal{B})} \\ &= \frac{P(\mathcal{N}|\mathcal{C}, \mathcal{X}, \mathcal{B}')P(\mathcal{B}')}{P(\mathcal{N}|\mathcal{C}, \mathcal{X}, \mathcal{B})P(\mathcal{B})}, \end{aligned} \quad (15)$$

where $P(\mathcal{N}|\mathcal{C}, \mathcal{X}, \mathcal{B})$ is the partial likelihood in Equation (4).

For $P(\mathcal{B})$, we select a multivariate Gaussian priors as mentioned earlier. Similar to what we did in Section 3.1, we can take the log and obtain the log of acceptance ratio as following:

$$\begin{aligned} &\log\left(\phi_d(\boldsymbol{\beta}'^{(c)}; \mathbf{0}, \sigma^2 I_P)\right) - \log\left(\phi_d(\boldsymbol{\beta}'^{(c)}; \mathbf{0}, \sigma^2 I_P)\right) \\ &+ \sum_{d:c^{(d)}=c} \left\{ \boldsymbol{\beta}'^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)}) - \log\left[\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}'^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j)\} \right] \right\} \\ &- \sum_{d:c^{(d)}=c} \left\{ \boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j^{(d)}) - \log\left[\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} x_{t^{(d)}}^{(c)}(i^{(d)}, j)\} \right] \right\}, \end{aligned} \quad (16)$$

where $\phi_d(\cdot; \mu, \Sigma)$ is the d -dimensional multivariate normal density. Then the log of acceptance ratio we have is:

$$\log(\text{Acceptance Probability}) = \min((16), 0) \quad (17)$$

To determine whether we accept the proposed update or not, we take the usual approach, by comparing the log of acceptance ratio we have to the log of a sample from $\text{uniform}(0,1)$.

3.4 Pseudocode

To implement the inference procedure outlined above, we provide a pseudocode for Markov Chain Monte Carlo (MCMC) sampling. Note that we use two loops, outer iteration and inner iteration, in order to avoid the label switching problem (Jasra et al., 2005), which is an issue caused by the nonidentifiability of the components under symmetric priors in Bayesian mixture modeling. When summarizing model results, we will only use the values from the last I^{th} outer loop because there is no label switching problem within the inner iteration.

Algorithm 5 MCMC($I, n_1, n_2, n_3, \delta_B$)

set initial values $\mathcal{C}^{(0)}$, $\mathcal{Z}^{(0)}$, and $\mathcal{B}^{(0)}$

for $i=1$ to I **do**

for $n=1$ to n_1 **do**

 fix $\mathcal{Z} = \mathcal{Z}^{(i-1)}$ and $\mathcal{B} = \mathcal{B}^{(i-1)}$

for $d=1$ to D **do**

 calculate $\mathbf{x}_{t^{(d)}}^{*(c)}(i^{(d)}, j)$ according to Section 2.3, for every $c = 1, \dots, C$

 calculate $p^{\mathcal{C}}|\mathbf{z}^{(d)}, \boldsymbol{\beta}^{(c^{(d)})} = (p_1, \dots, p_C)$, where $p_c = \exp(\text{Eq. (11) corresponding to } c)$

 draw $c^{(d)} \sim \text{multinomial}(p^{\mathcal{C}})$

end

end

for $n=1$ to n_2 **do**

 fix $\mathcal{C} = \mathcal{C}^{(i)}$ and $\mathcal{B} = \mathcal{B}^{(i-1)}$

for $d=1$ to D **do**

for $m=1$ to $M^{(d)}$ **do**

 calculate $p^{\mathcal{Z}}|\mathbf{c}^{(d)}, \boldsymbol{\beta}^{(c^{(d)})} = (p_1, \dots, p_K)$, where $p_k = \exp(\text{Eq. (13) corresponding to } k)$

 draw of $z_m^{(d)} \sim \text{multinomial}(p^{\mathcal{Z}})$

end

end

end

for $n=1$ to n_3 **do**

 fix $\mathcal{C} = \mathcal{C}^{(i)}$, $\mathcal{Z} = \mathcal{Z}^{(i)}$, and $\mathcal{B}^{(0)} = \text{last value } (n_3^{th}) \text{ of } \mathcal{B}^{(i-1)}$

 calculate $\mathcal{X} = \{\mathbf{x}_{t^{(d)}}^{*(c)}(i, j)\}_{d=1}^D$ according to Section 2.3, given fixed \mathcal{C}

for $c=1$ to C **do**

 draw $\boldsymbol{\beta}^{(c)}|\mathcal{C}, \mathcal{Z}, \mathcal{B}^{(n-1)}$ using M-H algorithm in Section 3.3

end

end

end

summarize the results using:

the last value of \mathcal{C} , the last value of \mathcal{Z} , and the last n_3 length chain of \mathcal{B}

4 Application: North Carolina email data

To see the applicability of the model, we used the North Carolina email data using two counties, Vance county and Dare county, which are the two counties whose email corpus cover the date of Hurricane Sandy (October 22, 2012 – November 2, 2012). Exploratory analysis revealed that Dare county experienced significant change in the pattern of email exchanges; specifically, during the emergency period, email interactions significantly less rely on previous history of interactions, compared to the normal period. On the other hand, Vance county did not experience any distinctive change, and the possible reason for the difference is the locations of two counties. Here we apply IPTM to both data to see the differences in detail, in terms of the interaction patterns and topics of the corpus. One thing to note here is that instead of using 4 different triadic covariates (2-send, 2-receive, sibling, cosibling), we added the four statistics and defined the new statistic ‘triangles’, in order to simply measure triangle effects in general.

4.1 Vance county email data

After treating multicast emails (those involving a single sender but multiple receivers) as multiple distinct emails, Vance county data contains 269 emails (only count the email with the number of words greater than 0) between 18 actors, including 620 vocabulary in total. We used $K = 20$ topics assuming symmetric Dirichlet prior with the concentration parameter $\alpha = 5$, and $C = 3$ interaction patterns assuming multinomial prior with parameter γ (coming from symmetric Dirichlet prior with the concentration parameter $\eta = 5$). For topic-word distributions, we assumed that ϕ follows

symmetric Dirichlet distribution with the concentration parameter $\delta = 5$. MCMC sampling was implemented based on the order and scheme illustrated in Section 3. We set the outer iteration number as $I = 100$, and inner iteration numbers as $n_1 = 10$, $n_2 = 10$, and $n_3 = 3500$, which took about 5.5 hours in total. In addition, after some experimentation, δ_B was set as 0.5, to ensure sufficient acceptance rate (IP1: 0.129, IP2: 0.576, IP3: 0.506). In our case, the average acceptance rate for β was 0.260. As demonstrated in Algorithm 5, the last value of \mathcal{C} , the last value of \mathcal{Z} , and the last n_3 length chain of \mathcal{B} were taken as the final posterior samples. Among the \mathcal{B} samples, 500 were discarded as a burn-in, and every 3rd sample was taken for thinning. After these post-processing, MCMC diagnostic plots are attached in APPENDIX C, as well as geweke test statistics. There are some evidence of slightly bad mixing, which could be overcome if we sacrifice computation time and increase the size of thinning or iterations.

Below are the summary of IP-topic-word assignments. Each interaction pattern is paired with (a) posterior estimates of dynamic network effects corresponding to the interaction pattern, (b) the top 3 topics most likely to be generated conditioned on the interaction pattern, and (c) the top 10 most likely words to have generated conditioned on the topic and interaction pattern.

	IP1 (127 emails)	IP2 (56 emails)	IP3 (86 emails)
intercept	-0.110 [-1.76, 1.84]	-0.189 [-2.04, 1.68]	0.014 [-1.89, 1.75]
send	2.345 [1.71, 3.15]	1.017 [-0.83, 2.86]	1.021 [-0.21, 2.22]
receive	1.530 [0.82, 2.31]	0.770 [-1.26, 2.31]	1.176 [-0.30, 2.53]
triangles	0.589 [0.34, 0.82]	1.418 [-0.04, 2.52]	0.047 [-1.32, 1.11]

Table 2: Summary of posterior estimates of $\beta^{(c)}$ for Vance county emails

First, Table 2 summarizes the posterior means and 95% credible interval for $\beta^{(c)}$ corresponding to each interaction patterns. Below are the several examples of the interpretation of estimates, in the context of point process framework in Section 2.1. The interpretation can be extended to any other interaction patterns and time intervals between the two emails.

- **(Intercept)** Assuming no history at all between the sender and receiver, the document is $\frac{e^{(0.014)}}{e^{(-0.110)}} \approx 1.132$ times more likely to be IP3 relative to IP1, $\frac{e^{(0.014)}}{e^{(-0.189)}} \approx 1.225$ times more likely to be IP3 relative to IP2, and $\frac{e^{(-0.110)}}{e^{(-0.189)}} \approx 1.082$ times more likely to be IP2 relative to IP1
- **(Send)** If i sends an email to j at time t , the likelihoods of i sends email of IP1 to j after 1 hour and 2 hours are multiplied by $e^{(2.345 \times e^{(-0.05)})} \approx 9.306$ and $e^{(2.345 \times e^{(-2 \times 0.05)})} \approx 8.347$, respectively.
- **(Receive)** If j sends an email to i at time t , the likelihoods of i sends email of IP1 to j after 30 minutes and 30 hours are multiplied by $e^{(1.530 \times e^{(-0.5 \times 0.05)})} \approx 4.447$ and $e^{(1.530 \times e^{(-30 \times 0.05)})} \approx 1.407$, respectively.
- **(Triangles)** If i sends/receives an email to/from k at time t , and k sends/receives an email to/from j at time $t + 1$ (i.e. after 1 hour from t), then i sends/receives email to/from j at time $t + 2$ (i.e. after 2 hour from t) at a slightly higher rate if IP1 (likelihood multiplied by $e^{(0.589 \times e^{(-1)} \times e^{(-1)})} \approx 1.083$).

By examining the estimates in Table 2 and their corresponding interpretation, it seems that there exist strong effects of dynamic network covariates. That is, whether the sender and receiver previously had dyadic or triangle interaction strongly increase the rate of their interactions. Moreover, to see the differences across the interaction patterns more clearly, we compared the posterior distribution using the boxplots in Figure 2 and it seems that there exists notable differences in dynamic network covariates across the interaction patterns. For example, IP1 has the highest send and receive effect, while its baseline intensity (i.e. intercept) or triangle effect is not as high as other interaction patterns. Later, multiple hypothesis testing could be applied in order to test the significance of the differences in $\beta^{(c)}$ across the C number of interaction patterns.




Figure 2: Posterior distribution of $\beta^{(c)}$ for Vance county emails

Next, we scrutinize the topic distributions corresponding to each interaction patterns in Figure 3. There is some distinctive differences in the topic distributions \mathcal{Z} , given the assignment of interaction patterns to the documents \mathcal{C} . Specifically, each interaction pattern has different topics as the topic with highest probability (IP1: topic 9, IP2: topic 2, IP3: topic 11).




Figure 3: Posterior distribution of \mathcal{Z} for Vance county emails

Furthermore, we look at the distribution of words given the topics, which corresponds to Algorithm 4 in the generative process. Since the topic-word distribution ϕ does not depend on the interaction patterns as previous cases, Table 3 lists top 10 topics with top 10 words that have the highest probability conditioned on the topic. In addition, this time we try to check the interaction pattern-word distribution by listing top 10 words that have the highest probability conditioned on the interaction pattern. It seems that the words are not significantly different, having several words like ‘director’, ‘phones’, ‘department’, ‘description’, or ‘henderson’ (county seat of Vance county) appeared repetitively across the most of the topics or interaction patterns. The word ‘will’ was ranked the top in most of the lists, probably because it was not deleted during the text mining process while other similar type of words like ‘am’, ‘is’, ‘are’, or ‘can’ are all removed.

Topic 9 (0.0946)		Topic 19 (0.0907)		Topic 2 (0.0869)		Topic 13 (0.0849)		Topic 8 (0.0671)	
will	0.0622	will	0.0487	will	0.709	will	0.0237	will	0.0419
suite	0.0184	director	0.0428	director	0.0370	street	0.0221	system	0.0279
henderson	0.0170	description	0.0310	department	0.0277	phones	0.0189	phones	0.0259
october	0.0170	message	0.0206	henderson	0.0200	october	0.0173	director	0.0200
system	0.0170	phones	0.0177	center	0.0154	fax	0.0158	center	0.0180
extension	0.0156	meeting	0.0162	phone	0.0154	suite	0.0158	emergency	0.0160
meeting	0.0156	fax	0.0147	october	0.0139	church	0.0142	october	0.0160
phone	0.0156	latest	0.0147	street	0.0139	phone	0.0142	attached	0.0140
electronic	0.0141	october	0.0147	church	0.0123	advised	0.0139	department	0.0140
heads	0.0141	street	0.0147	fax	0.0123	heads	0.0126	message	0.0140
Topic 15 (0.0628)		Topic 7 (0.0566)		Topic 3 (0.0435)		Topic 6 (0.0427)		Topic 16 (0.0400)	
will	0.0554	henderson	0.0378	heads	0.0308	department	0.0502	operations	0.0268
henderson	0.0256	will	0.0355	will	0.0277	will	0.0408	phone	0.0268
phones	0.0256	phone	0.0307	director	0.0246	director	0.0345	system	0.0268
phone	0.0192	church	0.0189	street	0.0215	phone	0.0251	department	0.0234
street	0.0192	coming	0.0189	directory	0.0185	electronic	0.0219	church	0.0201
suite	0.0192	phones	0.0189	phone	0.0185	week	0.0219	director	0.0201
cutting	0.0171	suite	0.0165	attached	0.0154	church	0.0157	message	0.0201
department	0.0171	training	0.0165	days	0.0154	description	0.057	phones	0.0201
electronic	0.0149	attached	0.0142	description	0.0154	development	0.0157	will	0.0201
rest	0.0149	description	0.0142	e-mail	0.0154	emergency	0.0157	directory	0.0167

Table 3: Summary of top 10 topics with top 10 words that have the highest probability conditioned on the topic

IP1 (0.4721)		IP2 (0.2082)		IP3 (0.3197)	
will	0.0505	will	0.0551	director	0.257
director	0.0201	director	0.0187	will	0.0216
phones	0.0198	department	0.0177	henderson	0.0199
henderson	0.0164	description	0.0167	operations	0.0193
phone	0.0164	phones	0.0152	street	0.0187
department	0.0159	phone	0.0141	emergency	0.0175
street	0.0156	henderson	0.0136	fax	0.0164
system	0.0148	street	0.0131	church	0.0140
october	0.0132	heads	0.0126	suite	0.0140
week	0.0119	meeting	0.0111	latest	0.0134

Table 4: Summary of top 10 words that have the highest probability conditioned on the interaction patterns

Although Vance county email data did not display distinctive idiosyncrasy across the interaction patterns and the topic-token assignments, it is not surprising because Vance county is a small county (land area: 253.52 sq. mi and population: 44,998), and our exploratory data analysis did not find any significant change in the email exchanges of department managers during the period of hurricane Sandy. Yet, it is definitely worthwhile to further look at this in terms of showing the applicability of interaction-partitioned topic model (IPTM), in case of email data. In the next section, we apply the methods for implementing an asymmetric Dirichlet prior in Wallach (2008) and Wallach et al. (2009), in hope of improving the model fitting and finding more interesting results in terms of interaction patterns-topics-words relationship.

5 Asymmetric Dirichlet prior over Θ (topic distribution)

Wallach et al. (2009) demonstrated that the typical implementations of topic models using symmetric Dirichlet priors with fixed concentration parameters often result in less practical results, which could be our case shown in Section 4.1. Commonly, the model fitting is improved by applying an asymmetric Dirichlet prior over the document–topic distributions (i.e. Θ), while an asymmetric prior over the topic–word distributions (i.e. Φ) provides no real benefit. Therefore, we assign an asymmetric Dirichlet prior over the interaction pattern–topic distributions, $\Theta = \{\theta^{(c)}\}_{c=1}^C$, where $\theta^{(c)}$ is drawn from $\text{Dir}(\alpha, \mathbf{m})$. While Wallach et al. (2009) illustrates two different methods, adding a hierarchy to Θ and optimizing the hyperparameters (α and \mathbf{m}), we choose to use hyperparameter optimization steps since it is computationally efficient and also sufficient to achieve the desired performance gains. Now, we assume \mathbf{m} to be nonuniform base measures (while α is still a fixed concentration parameter), and implement the hyperparameter optimization technique called “new fixed-point iterations using the Digamma recurrence relation” in Wallach (2008) based on Minka’s fixed-point iteration (Minka, 2000).

5.1 New fixed-point iterations using the Digamma recurrence relation

In this section, we summarize Chapter 2 of Wallach (2008) and its extension to our IPTM, to illustrate the basic steps and equations used for our optimization. Basically, we want to find the optimal hyperparameter $[\alpha\mathbf{m}]^*$ given the data \mathcal{D} such that the probability of the data given the hyperparameters $P(\mathcal{D}|\alpha\mathbf{m})$ is maximized at $[\alpha\mathbf{m}]^*$. After incorporating the interaction pattern, the evidence is now given by

$$P(\mathcal{D}|\alpha\mathbf{m}) = \prod_{d=1}^D \frac{\Gamma(\alpha)}{\Gamma(M_{c^{(d)}}^{\cdot} + \alpha)} \prod_{k=1}^K \frac{\Gamma(M_{c^{(d)}k}^{CK} + \alpha m_k)}{\Gamma(\alpha m_k)} \quad (18)$$

and is concave in $\alpha\mathbf{m}$, thus we will estimate $[\alpha\mathbf{m}]^*$ based on some pilot runs of MCMC.

First, the starting point is derived by Minka’s fixed-point iteration which takes the derivative of the lower bound $B([\alpha\mathbf{m}]^*)$ of $\log P(\mathcal{D}|\alpha\mathbf{m})$ with respect to $[\alpha m_k]^*$:

$$[\alpha m_k]^* = \alpha m_k \frac{\sum_{d=1}^D \Psi(M_{c^{(d)}k}^{CK} + \alpha m_k) - \Psi(\alpha m_k)}{\sum_{d=1}^D \Psi(M_{c^{(d)}}^{\cdot} + \alpha) - \Psi(\alpha)}, \quad (19)$$

where $\Psi(\cdot)$ is the first derivative of the log gamma function, known as the digamma function, and the quantity $M_{c^{(d)}k}^{CK}$ is the number of times that outcome k was observed across the documents of the interaction pattern $c^{(d)} \in \{1, \dots, C\}$. Moreover, the quantity $M_{c^{(d)}}^{\cdot} = \sum_{k=1}^K M_{c^{(d)}k}^{CK}$ is the total number of words across the documents of the interaction pattern $c^{(d)} \in \{1, \dots, C\}$. The value αm_k acts as an initial “pseudocount” for outcome k in all documents.

Next, Wallach’s new method rewrites the equation above using the notation $C_k(n) = \sum_{d=1}^D \delta(M_{c^{(d)}k}^{CK} - n)$ and $C.(n) = \sum_{d=1}^D \delta(M_{c^{(d)}}^{\cdot} - n)$:

$$[\alpha m_k]^* = \alpha m_k \frac{\sum_{n=1}^{\max_d M_{c^{(d)}k}^{CK}} C_k(n) [\Psi(n + \alpha m_k) - \Psi(\alpha m_k)]}{\sum_{n=1}^{\max_d M_{c^{(d)}}^{\cdot}} C.(n) [\Psi(n + \alpha) - \Psi(\alpha)]}. \quad (20)$$

Finally, applying the digamma recurrence relation (for any positive integer n)

$$\Psi(n + z) - \Psi(z) = \sum_{f=1}^n \frac{1}{f - 1 + z},$$

we substitute Equation (20) for below:

$$[\alpha m_k]^* = \alpha m_k \frac{\sum_{n=1}^{\max_d M_{c^{(d)}k}^{CK}} C_k(n) \sum_{f=1}^n \frac{1}{f - 1 + \alpha m_k}}{\sum_{n=1}^{\max_d M_{c^{(d)}}^{\cdot}} C.(n) \sum_{f=1}^n \frac{1}{f - 1 + \alpha}}. \quad (21)$$

This method is as accurate as Mika’s fixed-point iteration method, but it achieves computational efficiency since the digamma recurrence relation reduces the number of new calculations required for each successive n to one. Pseudocode for the complete fixed-point iteration is given in algorithm 2.2 of Wallach (2008).

5.2 Application to Vance county email data

The optimization procedure in Section 5.1 was applied to Vance county email data, using the settings $K = 10$ and $C = 3$ but assuming asymmetric base prior \mathbf{m} (and not fixing α as well). To see the advantage of optimizing α and \mathbf{m} in terms of model fitting, we computed the proportionality constants of $\log P(\mathcal{W}, \mathcal{Z} | \mathcal{C}, \delta, \mathbf{n}, \alpha, \mathbf{m})$ in Equation (13) of Section 3.2, and compared with those of the symmetric model. This time, we fix the parameters as $\alpha = 50/K = 50/10 = 5$ with $\mathbf{m} = (1, \dots, 1)^T$, $\delta = 0.01$, and $\eta = 50/C = 50/3 = 16.67$, following the common practice. The asymmetric model uses same δ and η , but α and \mathbf{m} are optimized for every outer iteration. Figure 4 below compares the MCMC results of symmetric and asymmetric models.

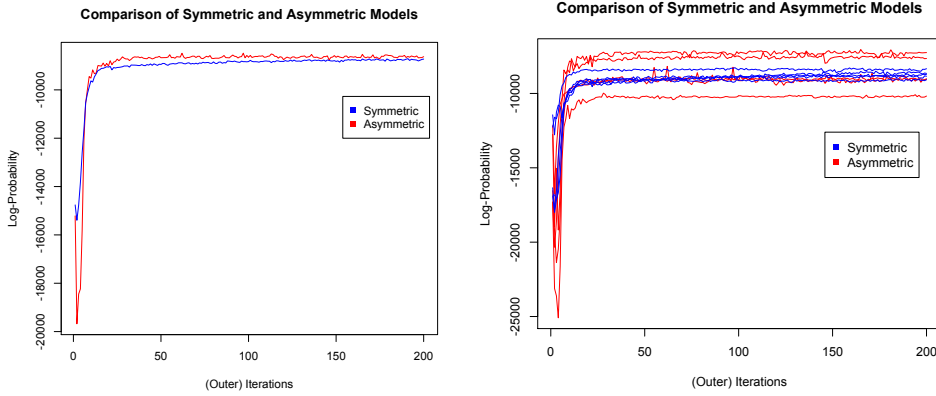


Figure 4: The proportionality constant of $\log P(\mathcal{W}, \mathcal{Z} | \mathcal{C}, \delta, \mathbf{n}, \alpha, \mathbf{m})$ for symmetric and asymmetric models, averaged over 5 Gibbs sampling runs (left) and individual runs (right)

Moreover, to measure the within-model consistency of interaction pattern and topic assignments (i.e. similarity between 1) two sets of interaction pattern assignments \mathcal{C} and \mathcal{C}' , and 2) two sets of topic assignments \mathcal{Z} and \mathcal{Z}'), we computed the distance between different partitions of the tokens into $K = 10$ topics, using variation of information (VI) (Meilă, 2003). Following Wallach et al. (2009), we calculated the average VI distance between all 10 unique pairs of topic assignments from the 5 Gibbs runs for the symmetric and asymmetric model, respectively. We also calculated the between-model VI distance, averaged over all 25 unique pairs of topic assignments for the symmetric and asymmetric model. Since VI measures variability, the higher the value of VI, the greater the dissimilarity between the two assignments.

$\text{VI}(\mathcal{C}, \mathcal{C}')$	Symmetric \mathbf{m}	Asymmetric \mathbf{m}
Symmetric \mathbf{m}	1.523 ± 0.108	1.590 ± 0.198
Asymmetric \mathbf{m}	-	1.717 ± 0.185

Table 5: Average VI distances of \mathcal{C} between 5 runs of symmetric and asymmetric model

$\text{VI}(\mathcal{Z}, \mathcal{Z}')$	Symmetric \mathbf{m}	Asymmetric \mathbf{m}
Symmetric \mathbf{m}	3.125 ± 0.212	3.341 ± 0.234
Asymmetric \mathbf{m}	-	3.135 ± 0.160

Table 6: Average VI distances of \mathcal{Z} between 5 runs of symmetric and asymmetric model

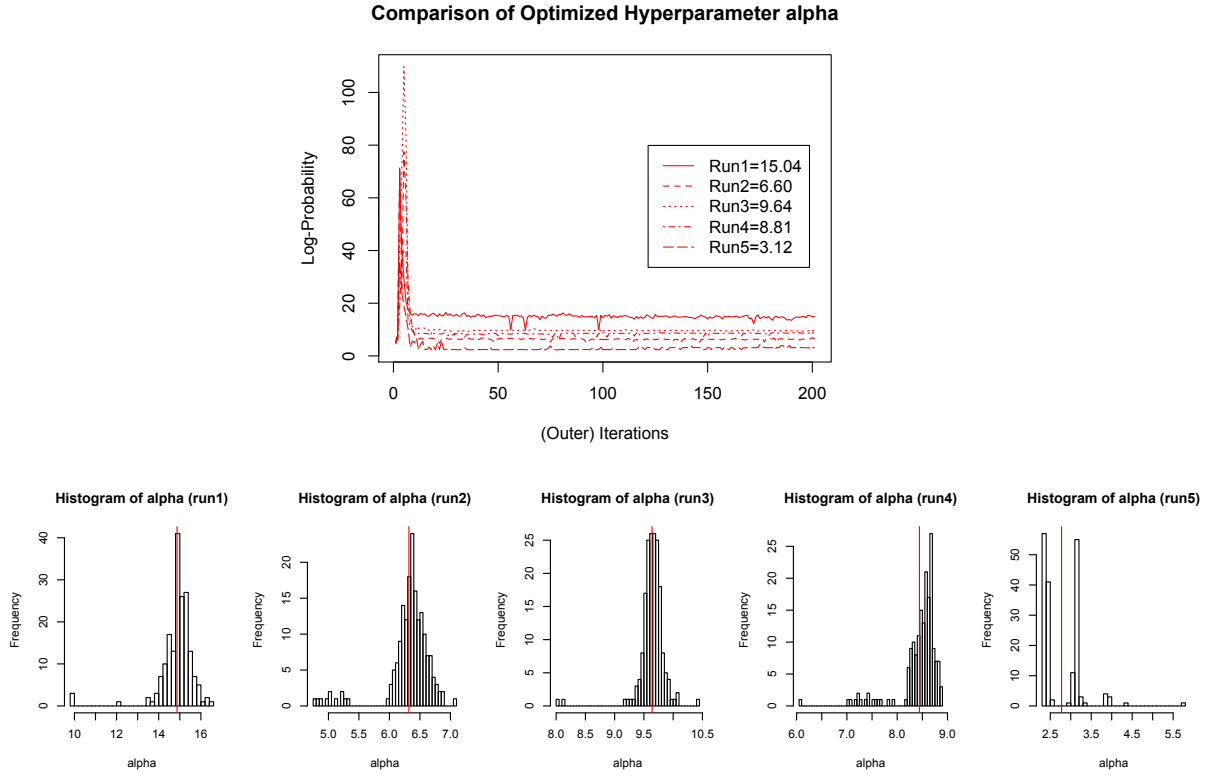


Figure 5: Plots and histograms of concentration parameter values α for 5 Gibbs sampling runs

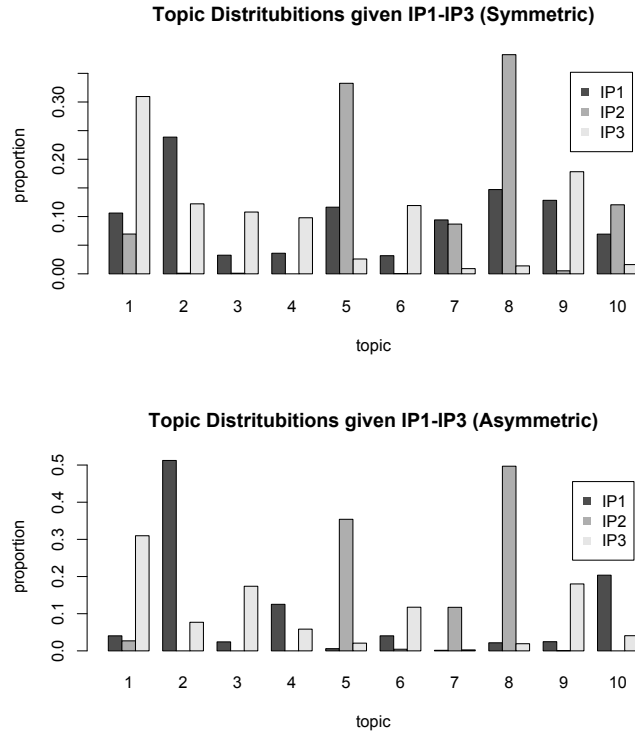


Figure 6: Posterior distribution of \mathcal{Z} for Vance county emails, using the Run5

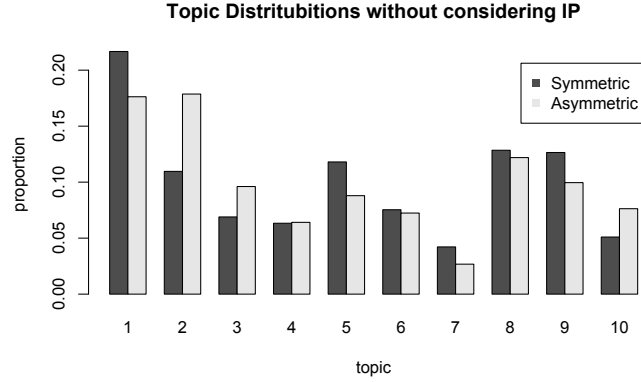


Figure 7: Posterior distribution of \mathcal{Z} for Vance county emails

Topic 1 (0.2167)	Topic 8 (0.1285)	Topic 9 (0.1265)	Topic 5 (0.1180)	Topic 2 (0.1096)
church center meeting heads office department development enp good pursuant	directory october jail provided technology training cutting rest attached day	suite message request cem records attachments hereto morning phase financial	will phones week phone system advised today latest tuesday monday	henderson street description director planning attached latest email wanted excel

Table 7: Summary of top 5 topics with top 10 words that have the highest probability conditioned on the topic (Symmetric)

Topic 2 (0.1787)	Topic 1 (0.1763)	Topic 8 (0.1219)	Topic 9 (0.0996)	Topic 3 (0.0961)
meesage electronic time request review sure meeting heads jail response	fax church suite center office development good email contacting technology	will department provided cutover training cutting day monday questions folks	street emergency phone communications cem board tomorrow spread sorted taps	operations description latest planning enp e-mail goldvance sealimprovedjordan phase wanted rural

Table 8: Summary of top 5 topics with top 10 words that have the highest probability conditioned on the topic (Asymmetric)

Symmetric	IP1 (54 emails)	IP2 (52 emails)	IP3 (163 emails)
intercept	0.341 [-0.66, 1.54]	1.160 [-0.11, 2.65]	1.168* [0.32, 2.38]
send	1.407 [-0.24, 2.98]	0.297 [-1.66, 2.27]	1.082* [0.53, 1.75]
receive	-0.621 [-2.56, 1.03]	-0.539 [-2.30, 1.02]	0.402 [-0.58, 1.22]
triangles	0.825 [-0.20, 2.05]	0.472 [-0.60, 1.42]	-0.407 [-1.74, 0.53]
Asymmetric	IP1 (81 emails)	IP2 (26 emails)	IP3 (162 emails)
intercept	1.316* [0.08, 2.58]	0.750 [-0.67, 2.34]	0.646 [-0.267, 1.57]
send	0.895 [-0.89, 2.64]	0.170 [-1.41, 1.98]	1.214* [0.73, 1.64]
receive	-0.242 [-2.18, 1.47]	-0.448 [-2.30, 1.21]	-0.349 [-1.27, 0.41]
triangles	1.028* [0.29, 1.97]	1.123 [-0.06, 2.34]	0.731 [0.20, 1.26]

Table 9: Posterior estimates of $\beta^{(c)}$: Symmetric (upper) and Asymmetric (lower)

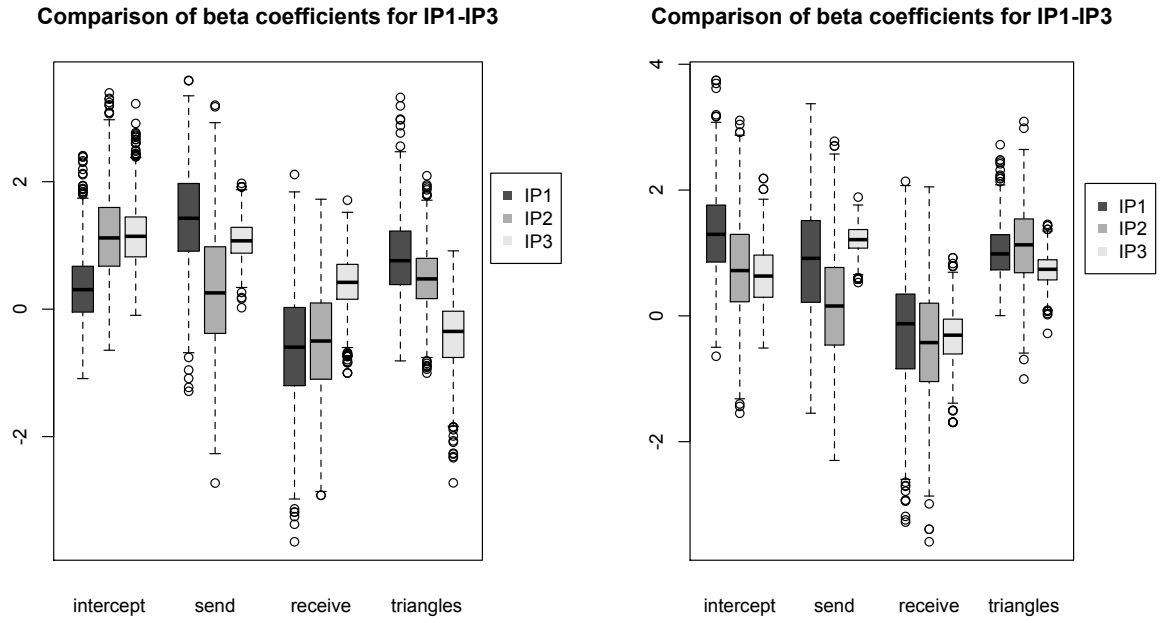


Figure 8: Posterior distribution of $\beta^{(c)}$ for Vance county emails (Asymmetric)

5.3 Application to Dare county email data

6 Asymmetric Dirichlet prior over Γ (IP distribution)

Since the introduction of interaction patterns is the novel aspect of our work, we extend previous Section 5 to our interaction pattern distribution Γ , in hope of significant improvement in model fitting. Therefore, we assign an asymmetric Dirichlet prior over the corpus-interaction pattern distributions, $\Gamma = \{\gamma\}$, where γ is drawn from $\text{Dir}(\eta, \mathbf{l})$. Again, we choose to use hyperparameter optimization steps instead of adding a new hierarchy. Now, we assume \mathbf{l} to be nonuniform base measures (while η is still a fixed concentration parameter), and implement the hyperparameter optimization technique called “new fixed-point iterations using the Digamma recurrence relation” in Wallach (2008) based on Minka’s fixed-point iteration (Minka, 2000).

6.1 New fixed-point iterations using the Digamma recurrence relation

In this section, we summarize Chapter 2 of Wallach (2008) to illustrate the basic steps and equations used for our optimization. Basically, we want to find the optimal hyperparameter $[\eta\mathbf{l}]^*$ given the data \mathcal{D} such that the probability of the data given the hyperparameters $P(\mathcal{D}|\eta\mathbf{l})$ is maximized at $[\eta\mathbf{l}]^*$. The evidence is given by

$$P(\mathcal{D}|\eta\mathbf{l}) = \frac{\Gamma(\eta)}{\Gamma(N. + \eta)} \prod_{c=1}^C \frac{\Gamma(N_c + \eta l_c)}{\Gamma(\eta l_c)} \quad (22)$$

and is concave in $\eta\mathbf{l}$, thus we will estimate $[\eta\mathbf{l}]^*$ based on some pilot runs of MCMC.

First, the starting point is derived by Minka’s fixed-point iteration which takes the derivative of the lower bound $B([\eta\mathbf{l}]^*)$ of $\log P(\mathcal{D}|\eta\mathbf{l})$ with respect to $[\eta l_c]^*$:

$$[\eta l_c]^* = \eta l_c \frac{\Psi(N_c + \eta l_c) - \Psi(\eta l_c)}{\Psi(N. + \eta) - \Psi(\eta)}, \quad (23)$$

where $\Psi(\cdot)$ is the first derivative of the log gamma function, known as the digamma function, and the quantity N_c is the number of times that interaction pattern c was observed in the corpus D . Moreover, the quantity $N. = \sum_{c=1}^C N_c$ is the total number of document in the corpus D . The value ηl_c acts as an initial “pseudocount” for outcome c in all documents.

Applying the digamma recurrence relation (for any positive integer n)

$$\Psi(n + z) - \Psi(z) = \sum_{f=1}^n \frac{1}{f - 1 + z},$$

we substitute Equation (20) for below:

$$[\eta l_c]^* = \eta l_c \frac{\sum_{f=1}^{N_c} \frac{1}{f - 1 + \eta l_c}}{\sum_{f=1}^{N.} \frac{1}{f - 1 + \eta}}. \quad (24)$$

APPENDIX

APPENDIX A: Deriving the sampling equations for IPTM

$$\begin{aligned}
& P(\Phi, \Theta, \mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2) \\
&= P(\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \Phi, \Theta, \mathcal{X}, \gamma, \boldsymbol{\eta}, \sigma^2) P(\Phi, \Theta | \delta, \mathbf{n}, \alpha, \mathbf{m}) \\
&= P(\mathcal{W} | \mathcal{Z}, \Phi) P(\mathcal{Z} | \Theta) P(\mathcal{N} | \mathcal{C}, \mathcal{B}, \mathcal{X}) P(\mathcal{B} | \mathcal{C}, \sigma^2) P(\Phi | \delta, \mathbf{n}) P(\Theta | \mathcal{C}, \alpha, \mathbf{m}) P(\mathcal{C} | \gamma) P(\gamma | \boldsymbol{\eta}) \\
&= \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} P(w_m^{(d)} | \phi_{z_m^{(d)}}) \right] \times \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} P(z_m^{(d)} | \boldsymbol{\theta}^{(c)}) \right] \times \left[\prod_{d=1}^D P(\mathbf{N}^{(d)}(t^{(d)}) | c^{(d)}, \mathbf{x}^{(c^{(d)})}(t^{(d)}), \boldsymbol{\beta}^{(c)}) \right] \\
&\quad \times \left[\prod_{c=1}^C P(\boldsymbol{\beta}^{(c)} | \sigma^2) \right] \times \left[\prod_{k=1}^K P(\phi^{(k)} | \delta, \mathbf{n}) \right] \times \left[\prod_{c=1}^C P(\boldsymbol{\theta}^{(c)} | \alpha, \mathbf{m}) \right] \times \left[\prod_{d=1}^D P(c^{(d)} | \gamma) \right] \times P(\gamma | \boldsymbol{\eta})
\end{aligned} \tag{25}$$

Since $P(\boldsymbol{\beta}^{(c)} | \sigma^2)$ is $\text{Normal}(\mathbf{0}, \sigma^2)$ and $P(\gamma | \boldsymbol{\eta})$ is $\text{Dirichlet}(\boldsymbol{\eta})$, we can drop the two terms out and further rewrite the equation (20) as below:

$$\begin{aligned}
& \propto \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} P(w_m^{(d)} | \phi_{z_m^{(d)}}) \right] \times \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} P(z_m^{(d)} | \boldsymbol{\theta}^{(c)}) \right] \times \left[\prod_{d=1}^D P(\mathbf{N}^{(d)}(t^{(d)}) | c^{(d)}, \mathbf{x}^{(c^{(d)})}(t^{(d)}), \boldsymbol{\beta}^{(c)}) \right] \\
&\quad \times \left[\prod_{k=1}^K P(\phi^{(k)} | \delta, \mathbf{n}) \right] \times \left[\prod_{c=1}^C P(\boldsymbol{\theta}^{(c)} | \alpha, \mathbf{m}) \right] \times \left[\prod_{d=1}^D P(c^{(d)} | \gamma) \right] \\
&= \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} \phi_{w_m^{(d)} z_m^{(d)}} \right] \times \left[\prod_{d=1}^D \prod_{m=1}^{M^{(d)}} \boldsymbol{\theta}_{z_m^{(d)}}^{(c)} \right] \times \left[\prod_{d=1}^D \frac{\exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c^{(d)})}(i^{(d)}, j^{(d)})\}}{\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c^{(d)})}(i^{(d)}, j)\}} \right] \\
&\quad \times \left[\prod_{k=1}^K \left(\frac{\Gamma(\sum_{w=1}^W \delta n_w)}{\prod_{w=1}^W \Gamma(\delta n_w)} \prod_{w=1}^W \phi_{wk}^{\delta n_w - 1} \right) \right] \times \left[\prod_{c=1}^C \left(\frac{\Gamma(\sum_{k=1}^K \alpha m_k)}{\prod_{k=1}^K \Gamma(\alpha m_k)} \prod_{k=1}^K (\boldsymbol{\theta}_k^{(c)})^{\alpha m_k - 1} \right) \right] \times \left[\prod_{d=1}^D \gamma_c^{I(c^{(d)}=c)} \right] \\
&= \left[\frac{\Gamma(\sum_{w=1}^W \delta n_w)}{\prod_{w=1}^W \Gamma(\delta n_w)} \right]^K \times \left[\frac{\Gamma(\sum_{w=1}^W \delta n_w)}{\prod_{w=1}^W \Gamma(\delta n_w)} \right]^C \times \left[\prod_{d=1}^D \frac{\exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c^{(d)})}(i^{(d)}, j^{(d)})\}}{\sum_{j \in \mathcal{A}^{(c)}} \exp\{\boldsymbol{\beta}^{(c)T} \mathbf{x}_{t^{(d)}}^{(c^{(d)})}(i^{(d)}, j)\}} \right] \\
&\quad \times \left[\prod_{d=1}^D \gamma_{c^{(d)}} \right] \times \left[\prod_{k=1}^K \prod_{w=1}^W \phi_{wk}^{M_{wk}^{WK} + \delta n_w - 1} \right] \times \left[\prod_{c=1}^C \prod_{k=1}^K (\boldsymbol{\theta}_k^{(c)})^{M_{ck}^{CK} + \alpha m_k - 1} \right]
\end{aligned} \tag{26}$$

where M_{wk}^{WK} is the number of times the w^{th} word in the vocabulary is assigned to topic k , and M_{ck}^{CK} is the number of times topic k shows up given the interaction pattern c . By looking at the forms of the terms involving Θ and Φ in Equation (21), we integrate out the random variables Θ and Φ , making use of the fact that the Dirichlet distribution is a conjugate prior of multinomial distribution.

Applying the well-known formula $\int \prod_{m=1}^M [x_m^{k_m-1} dx_m] = \frac{\prod_{m=1}^M \Gamma(k_m)}{\Gamma(\sum_{m=1}^M k_m)}$ to (22), we have:

$$\begin{aligned}
& P(\mathcal{W}, \mathcal{Z}, \mathcal{C}, \mathcal{B}, \mathcal{N} | \mathcal{X}, \delta, \mathbf{n}, \alpha, \mathbf{m}, \gamma, \boldsymbol{\eta}, \sigma^2) \\
&= \text{Const.} \int_{\Theta} \int_{\Phi} \left[\prod_{k=1}^K \prod_{w=1}^W \phi_{wk}^{M_{wk}^{WK} + \delta n_w - 1} \right] \left[\prod_{c=1}^C \prod_{k=1}^K (\boldsymbol{\theta}_k^{(c)})^{M_{ck}^{CK} + \alpha m_k - 1} \right] d\Phi d\Theta \\
&= \text{Const.} \left[\prod_{k=1}^K \int_{\phi_{:k}} \prod_{w=1}^W \phi_{wk}^{M_{wk}^{WK} + \delta n_w - 1} d\phi_{:k} \right] \times \left[\prod_{c=1}^C \int_{\theta_{:c}} \prod_{k=1}^K (\boldsymbol{\theta}_k^{(c)})^{M_{ck}^{CK} + \alpha m_k - 1} d\theta_{:c} \right] \\
&= \text{Const.} \left[\prod_{k=1}^K \frac{\prod_{w=1}^W \Gamma(M_{wk}^{WK} + \delta n_w)}{\Gamma(\sum_{w=1}^W M_{wk}^{WK} + \delta)} \right] \times \left[\prod_{c=1}^C \frac{\prod_{k=1}^K \Gamma(M_{ck}^{CK} + \alpha m_k)}{\Gamma(\sum_{k=1}^K M_{ck}^{CK} + \alpha)} \right].
\end{aligned} \tag{27}$$

APPENDIX B: Computing conditional probability

$$\begin{aligned}
& P(\mathbf{w}^{(d)}, \mathbf{z}^{(d)} | c^{(d)} = c, \mathcal{W}_{\setminus d}, \mathcal{Z}_{\setminus d}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m}) \\
& \propto \prod_{m=1}^{M^{(d)}} P(z_m^{(d)} = k, w_m^{(d)} = w | c^{(d)} = c, \mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m})
\end{aligned} \tag{28}$$

To obtain the Gibbs sampling equation, we need to obtain an expression for $P(z_m^{(d)} = k, w_m^{(d)} = w, c^{(d)} = c | \mathcal{W}_{\setminus d}, \mathcal{Z}_{\setminus d}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m})$. From Bayes' theorem and Gamma identity $\Gamma(k+1) = k\Gamma(k)$,

$$\begin{aligned}
& P(z_m^{(d)} = k, w_m^{(d)} = w, c^{(d)} = c | \mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m}) \\
& \propto \frac{P(\mathcal{W}, \mathcal{Z}, \mathcal{C} | \delta, \mathbf{n}, \alpha, \mathbf{m})}{P(\mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}_{\setminus d} | \delta, \mathbf{n}, \alpha, \mathbf{m})} \\
& \propto \frac{\prod_{k=1}^K \frac{\prod_{w=1}^W \Gamma(M_{wk}^{WK} + \delta n_w)}{\Gamma(\sum_{w=1}^W M_{wk}^{WK} + \delta)} \times \prod_{c=1}^C \frac{\prod_{k=1}^K \Gamma(M_{ck}^{CK} + \alpha m_k)}{\Gamma(\sum_{k=1}^K M_{ck}^{CK} + \alpha)}}{\prod_{k=1}^K \frac{\prod_{w=1}^W \Gamma(M_{wk, \setminus d, m}^{WK} + \delta n_w)}{\Gamma(\sum_{w=1}^W M_{wk, \setminus d, m}^{WK} + \delta)} \times \prod_{c=1}^C \frac{\prod_{k=1}^K \Gamma(M_{ck, \setminus d, m}^{CK} + \alpha m_k)}{\Gamma(\sum_{k=1}^K M_{ck, \setminus d, m}^{CK} + \alpha)}} \\
& \propto \frac{M_{wk, \setminus d, m}^{WK} + \delta n_w}{\sum_{w=1}^W M_{wk, \setminus d, m}^{WK} + \delta} \times \frac{M_{ck, \setminus d, m}^{CK} + \alpha m_k}{\sum_{k=1}^K M_{ck, \setminus d, m}^{CK} + \alpha}
\end{aligned} \tag{29}$$

Then, the conditional probability that a novel word generated in the document of interaction pattern $c^{(d)} = c$ would be assigned to topic $z_m^{(d)} = k$ is obtained by:

$$\begin{aligned}
& P(z_m^{(d)} = k | w_m^{(d)} = w, c^{(d)} = c, \mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m}) \\
& \propto \frac{M_{ck, \setminus d, m}^{CK} + \alpha m_k}{\sum_{k=1}^K M_{ck, \setminus d, m}^{CK} + \alpha}
\end{aligned} \tag{30}$$

In addition, the conditional probability that a new word generated in the document would be $w_m^{(d)} = w$, given that it is generated from topic $z_m^{(d)} = k$ is obtained by:

$$\begin{aligned}
& P(w_m^{(d)} = w | z_m^{(d)} = k, c^{(d)} = c, \mathcal{W}_{\setminus d, m}, \mathcal{Z}_{\setminus d, m}, \mathcal{C}_{\setminus d}, \delta, \mathbf{n}, \alpha, \mathbf{m}) \\
& \propto \frac{M_{wk, \setminus d, m}^{WK} + \delta n_w}{\sum_{w=1}^W M_{wk, \setminus d, m}^{WK} + \delta}
\end{aligned} \tag{31}$$

APPENDIX C: MCMC Diagnostics for Vance county emails




Figure 9: Plot of empirical entropy to check the distribution of IP assignments \mathcal{C}




Figure 10: Traceplots and density plots of $\beta^{(1)}$

```
> geweke.diag(mcmc)

Fraction in 1st window = 0.1
Fraction in 2nd window = 0.5

var1    var2    var3    var4
-3.7668  0.0945  0.4176 -0.2620
```

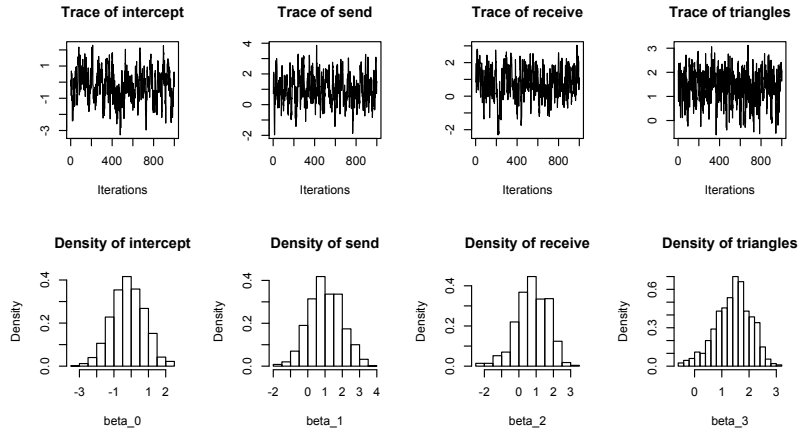


Figure 11: Traceplots and density plots of $\beta^{(2)}$

```
> geweke.diag(mcmc)

Fraction in 1st window = 0.1
Fraction in 2nd window = 0.5

var1    var2    var3    var4
0.02362 -0.22385 0.61747 -0.05824
```

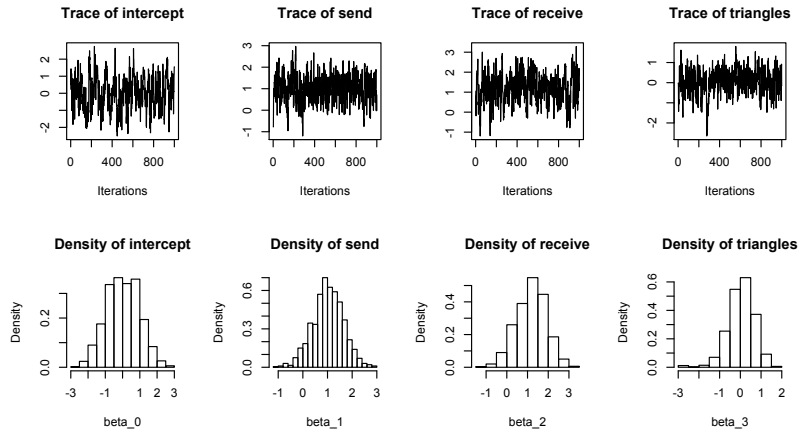


Figure 12: Traceplots and density plots of $\beta^{(3)}$

```
> geweke.diag(mcmc)

Fraction in 1st window = 0.1
Fraction in 2nd window = 0.5

var1    var2    var3    var4
0.5025  0.5471 -0.7871 -0.2434
```

References

- Blei, D. M., Ng, A. Y., and Jordan, M. I. (2003). Latent dirichlet allocation. *J. Mach. Learn. Res.*, 3:993–1022.
- Cox, D. R. (1992). Regression models and life-tables. In *Breakthroughs in statistics*, pages 527–541. Springer.
- Jasra, A., Holmes, C., and Stephens, D. (2005). Markov chain monte carlo methods and the label switching problem in bayesian mixture modeling. *Statistical Science*, pages 50–67.
- Meilă, M. (2003). Comparing clusterings by the variation of information. In *Learning theory and kernel machines*, pages 173–187. Springer.
- Minka, T. (2000). Estimating a dirichlet distribution.
- Perry, P. O. and Wolfe, P. J. (2013). Point process modelling for directed interaction networks. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 75(5):821–849.
- Rosen-Zvi, M., Griffiths, T., Steyvers, M., and Smyth, P. (2004). The author-topic model for authors and documents. In *Proceedings of the 20th conference on Uncertainty in artificial intelligence*, pages 487–494. AUAI Press.
- Vu, D. Q., Hunter, D., Smyth, P., and Asuncion, A. U. (2011). Continuous-time regression models for longitudinal networks. In Shawe-Taylor, J., Zemel, R., Bartlett, P., Pereira, F., and Weinberger, K., editors, *Advances in Neural Information Processing Systems 24*, pages 2492–2500. Curran Associates, Inc.
- Wallach, H. M. (2008). *Structured topic models for language*. PhD thesis, University of Cambridge.
- Wallach, H. M., Mimno, D. M., and McCallum, A. (2009). Rethinking lda: Why priors matter. In *Advances in neural information processing systems*, pages 1973–1981.